



# Kubernetes Introduction

DevOps Course

# Agenda

- ✦ Introduction
- ✦ High Level Architecture
- ✦ Kubernetes Objects and Workloads
- ✦ Demo

# What is Kubernetes?

- ✦ Is an open-source system for automating deployment, scaling, and management of containerized applications
- ✦ Kubernetes abstracts away the hardware infrastructure and exposes your whole datacenter as a single enormous computational resource.



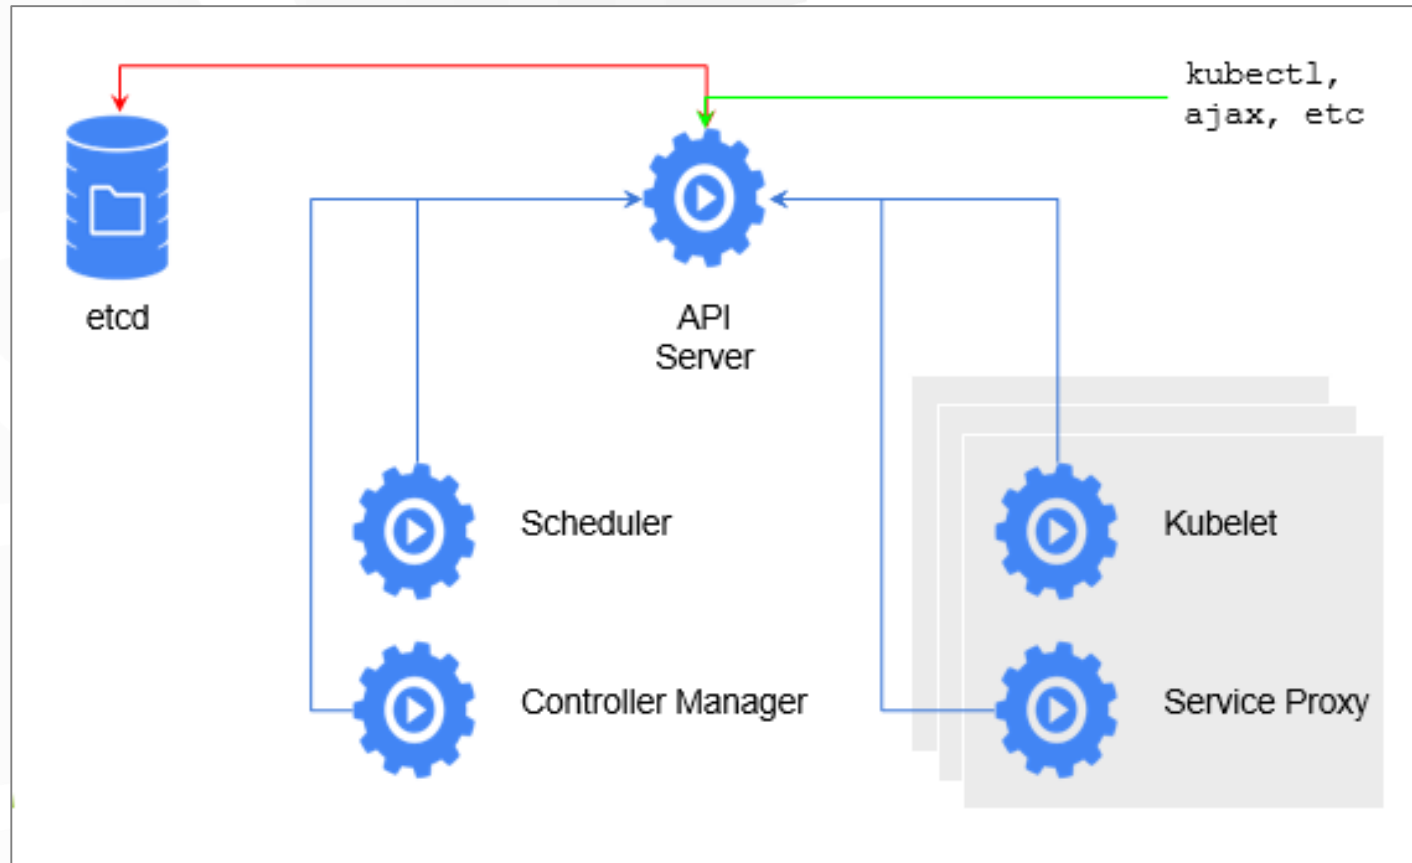
# Why Kubernetes?

- ✦ Moving from monolithic apps to microservices
- ✦ Providing a consistent environment to applications
- ✦ Moving to continuous delivery: DevOps and NoOps
- ✦ Containers need management
- ✦ Runs on any cloud & bare metal

# Main Features

- ✦ Monitoring
- ✦ Health Checking
- ✦ Horizontal Auto Scaling
- ✦ Service Discovery
- ✦ Rolling Deployment and Rollback
- ✦ Logging
- ✦ Load Balancing

# High Level Architecture



# etcd

- ✦ It stores the configuration information which can be used by each of the nodes in the cluster.
- ✦ It is a high availability key value store that can be distributed among multiple nodes.
- ✦ It is accessible only by Kubernetes API server as it may have some sensitive information.
- ✦ It is a distributed key value Store which is accessible to all.

# API Server

- ✦ Kubernetes is an API server which provides all the operation on cluster using the API
- ✦ API server implements an interface, which means different tools and libraries can readily communicate with it.
- ✦ Kubeconfig is a package along with the server side tools that can be used for communication. It exposes Kubernetes API.



# Controller Manager

- ✦ This component is responsible for most of the controllers that regulate the state of the cluster and perform tasks.
- ✦ In general, it can be considered as a daemon which runs in a nonterminating loop and is responsible for collecting and sending information to the API server.
- ✦ It works toward getting the shared state of the cluster and then makes changes to bring the current status of the server to the desired state.
- ✦ The controller manager runs different kinds of controllers to handle nodes, endpoints, etc.

# Scheduler

- ✦ It is a service in master responsible for distributing the workload
- ✦ It is responsible for tracking utilization of working load on cluster nodes and then placing the workload on which resources are available and accept the workload.
- ✦ The scheduler is responsible for workload utilization and allocating pod to new node.


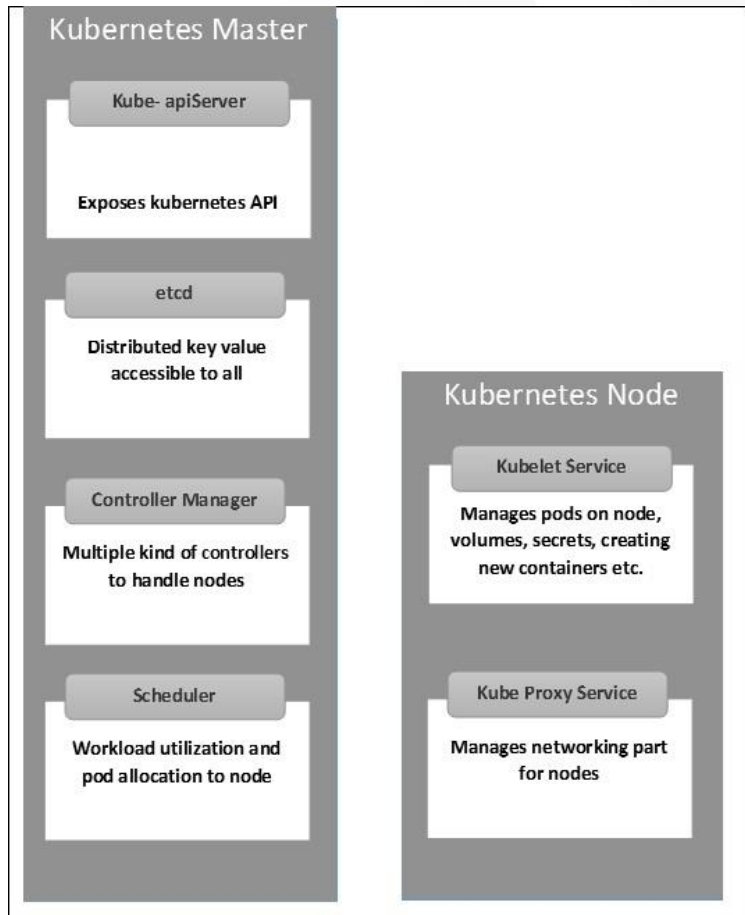
# Kubelet Service

- ✦ This is a small service in each node responsible for relaying information to and from control plane service
- ✦ It interacts with etcd store to read configuration details and write values.
- ✦ This communicates with the master component to receive commands and work.
- ✦ The kubelet process then assumes responsibility for maintaining the state of work and the node server.


# Kubernetes Proxy Service

- ✦ This is a proxy service which runs on each node and helps in making services available to the external host.
- ✦ It helps in forwarding the request to correct containers and is capable of performing primitive load balancing.
- ✦ It makes sure that the networking environment is predictable and accessible and at the same time it is isolated as well.
- ✦ It manages pods on node, volumes, secrets, creating new containers' health checkup, etc.

# Master and Node Structure



# Kubernetes Objects and Workloads



# Pods

- ✦ A pod is the most basic unit that Kubernetes deals with.
- ✦ Containers themselves are not assigned to hosts. Instead, one or more tightly coupled containers are encapsulated in an object called a pod.
- ✦ A pod generally represents one or more containers that should be controlled as a single application.
- ✦ users are encouraged to work with higher level objects that use pods or pod templates as base components.

# Replication Controllers

- ✦ Is an object that defines a pod template and control parameters to scale identical replicas of a pod horizontally by increasing or decreasing the number of running copies.
- ✦ The replication controller is responsible for ensuring that the number of pods deployed in the cluster matches the number of pods in its configuration.
- ✦ If the number of replicas in a controller's configuration changes, the controller either starts up or kills containers to match the desired number.



# Replication Sets

- ✦ are an iteration on the replication controller design with greater flexibility in how the controller identifies the pods it is meant to manage.
- ✦ Replication sets are meant to be used inside of additional, higher level units that provide that functionality.
- ✦ Replication sets are not able to do rolling updates to cycle backends to a new version like replication controllers can.

# Deployments

- ✦ Deployments are upgraded and higher version of replication controller. They manage the deployment of replica sets which is also an upgraded version of the replication controller.
- ✦ They have the capability to update the replica set and are also capable of rolling back to the previous version.
- ✦ Deployments are a high level object designed to ease the life cycle management of replicated pods.

# Jobs and Cron Jobs

- ✦ Provide a more task-based workflow where the running containers are expected to exit successfully after some time once they have completed their work.
- ✦ Jobs are useful if you need to perform one-off or batch processing instead of running a continuous service.
- ✦ Cron jobs can be used to schedule a job to execute in the future or on a regular, reoccurring basis.

# Services

- ✦ Is a component that acts as a basic internal load balancer and ambassador for pods.
- ✦ A service groups together logical collections of pods that perform the same function to present them as a single entity.
- ✦ Any time you need to provide access to one or more pods to another application or to external consumers, you should to configure a service.
- ✦ Although services, by default, are only available using an internally routable IP address, they can be made available outside of the cluster by choosing one of several strategies.

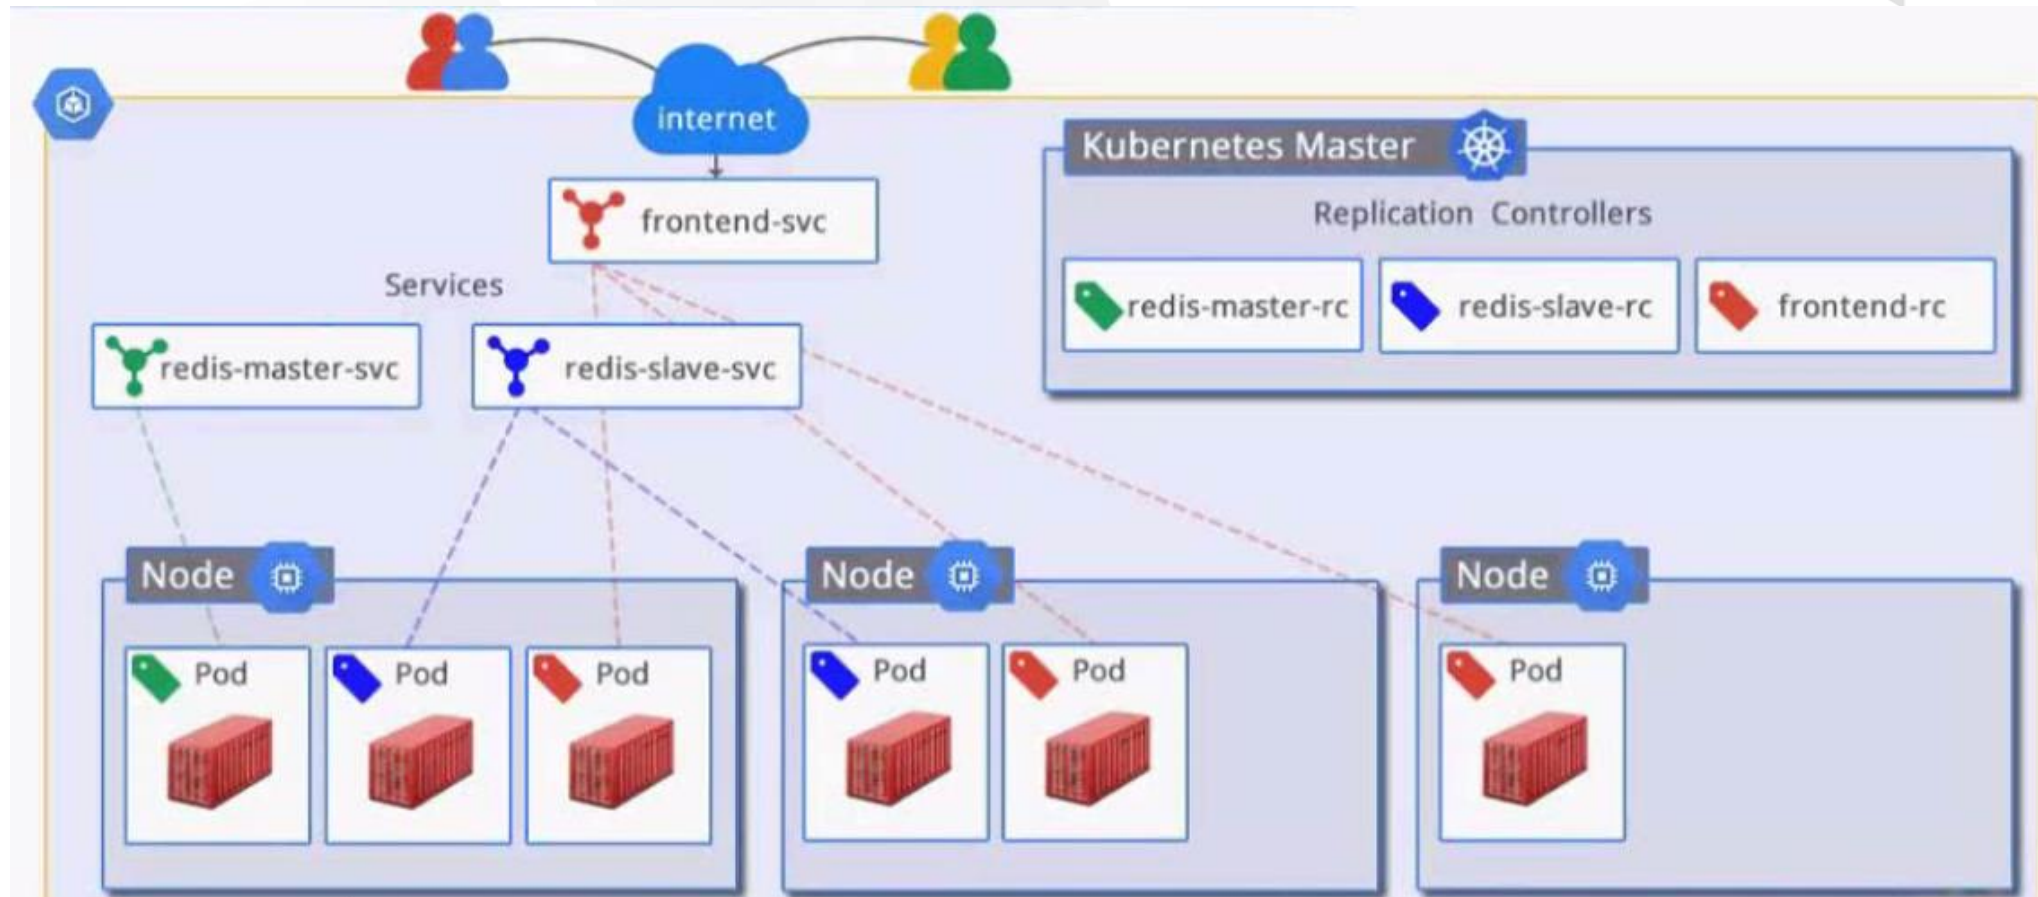
# Volumes

- ✦ Reliably sharing data and guaranteeing its availability between container restarts is a challenge in many containerized environments.
- ✦ Kubernetes uses its own volumes abstraction that allows data to be shared by all containers within a pod and remain available until the pod is terminated.
- ✦ Persistent volumes are a mechanism for abstracting more robust storage that is not tied to the pod life cycle.

# Labels and Annotations

- ✦ A label is a semantic tag that can be attached to Kubernetes objects to mark them as a part of a group.
- ✦ Services use labels to understand the backend pods they should route requests to.
- ✦ Annotations are a similar mechanism that allows you to attach arbitrary key-value information to an object.
- ✦ Annotations are a way of adding rich metadata to an object that is not helpful for selection purposes.

# Kubernetes Architecture



# Questions





# Introduction to Kubernetes

## Demo

